


Lecture 2

Introduction to Concurrency


Running a Sequential Program

- Executable
Machine instructions to be performed
- Program counter
Next instruction to be executed
- Stack
Current variable definitions
- Heap
Dynamically allocated data structures
- Control flow
Sequence of instructions performed during an execution



Java Memory Model

- Stack
 - Local variables
 - Method parameters
- Heap
 - Objects!
 - Every call to `new` allocates space on heap
- Class-typed variables reference heap or null



More on Main Memory (MM)

- Naively, MM is a table:
 - Each address can store a value
 - Each address refers to one memory location (no copies)

Address	Value
0000	'a'
0001	37
0002	NULL

- In reality, several copies of a given address are possible
 - Caches
 - Registers
 - ...
- Why? *Performance*
 - Higher-speed memory is more expensive
 - Copying frequently used data into high-speed memory (register, cache) improves performance while containing cost

Concurrent Programs

- Multiple control flows!
- Programs with multiple control flows can be
 - Concurrent
 - Parallel
 - Distributed
- Control flows are either
 - Processes
 - Threads

Concurrent vs. Parallel vs. Distributed

- Concurrent

of control flows unrelated to # of physical processors

- Parallel

of control flows \leq # of physical processors; each flow has its own processor

- Distributed


Multiple machines connected via network

Processes vs. Threads


- Processes
 - Possess own heap
 - Communicate via *IPC* (= inter-process communication) mechanisms
 - Sockets
 - Message passing
 - Etc.
- Threads
 - Contained within processes
 - Possess own stack, program counter
 - Share heap with other threads in same process
 - Communicate via shared memory
- Historically
 - Process management handled by operating system
 - Processes were single-threaded

(Single-Threaded) Processes


Process 1




IPC



Process 2



Multi-threaded Process




Running a Multi-Process / Multi-Threaded Application

- Execution requires processor
- Running a thread requires using a processor
- What decides which thread gets which processor?
 - Scheduler (part of operating system)!
 - Scheduling policy decides which threads run when
 - Pre-emptive schedulers can interrupt one thread and let another run on a given processor
 - Interrupted thread is “suspended”: its stack, program counter are saved so that thread can be re-activated later
 - Stack, program of new thread are loaded and new thread activated
 - This is called a *context switch*

Threads, Processes and Processors

- Do processes run on single machine? Yes
- Do processes run on a single processor? Not necessarily
 - Different threads can run on different processors
 - Scheduler makes this decision
- Do threads run on a single processor?
 - Usually
 - Some schedulers support *thread migration* (why?)

A Reference Model for Distributed / Parallel / Concurrent Programs



Language Support for Concurrency

- Many languages support concurrency!
C, C++, C#, OCaml, Java, Scala, Erlang, Python, ...
- Traditionally: process / thread management handled via system calls to operating system
 - Not part of core language (e.g. C)
 - Platform-specific, non-portable, since different OS's have different mechanisms
- Modern languages (e.g. Java, Scala, Erlang) include mechanisms for thread management directly

Java Concurrency

- Support for multi-threading, processes
 - Process = running instance of Java Virtual Machine
 - Objects live on heap, can be shared by threads in same process
- Every Java program has at least one thread: `main`
- This course: focus is on thread programming

Java Threads Are Objects

- Object class is `Thread`, which is part of `java.lang` package (automatically imported!)
- Thread objects include:
 - `public void run()` executed when thread is launched
 - `public void start()` to launch the thread
 - Other methods that we will study later
 - Constructors, of which more later, but here are two:
 - `Thread()` create a thread
 - `Thread(String name)` create a thread with the given name

Thread Creation in Java

- Create an object `t` in class `Thread` with desired functionality in `run()` method
- Invoke `t.start()`
- This starts a thread that runs the `t.run()` method!

“Desired Functionality in `run ()`”?

- Two approaches
 - Subclassing from `Thread`
 - Implementing `Runnable` interface
 - In the former: override `run ()`
 - In the second
 - Define a class implementing the `Runnable` interface
 - Use relevant constructor in `Thread` on objects in this class
- ```
Thread (Runnable target)
Thread (Runnable target, String name)
```

# Thread Implementation via Subclassing (Inheritance)

```
public class HelloWorldThread extends Thread {
 public void run() {
 System.out.println ("Thread says Hello World!");
 }
}
```

New class HelloWorldThread is introduced

- Extends Thread class
- Uses overriding to redefine `run ( )` method to do what we want

# Thread Implementation via Runnable

```
public class HelloWorldRunnable implements Runnable {
 public void run() {
 System.out.println ("Runnable says Hello World!");
 }
}
```

- Runnable is an interface in java.lang containing only:  
`public void run()`
- This class implements Runnable by providing each object with a `run()` method
- Constructor for Thread class can now be called with objects in this class

# Thread Creation

```
Thread h1 = new HelloWorldThread ();
Thread h2 = new Thread (new HelloWorldRunnable ());
h1.start();
h2.start();
```

- h1 is thread object created from subclass of Thread
- h2 is thread object created from Runnable object
- Output is two instances of “Hello World!”

# Subclassing or Runnable?

|      | Subclassing                                                                                                                                                                  | Runnable                                                                                                                                                                                                 |
|------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| PROS | <ul style="list-style-type: none"><li>• Easy access to Thread methods when implementing <code>run()</code></li><li>• No need for creating intermediate object</li></ul>      | <ul style="list-style-type: none"><li>• Can inherit from another class besides Thread when creating Runnable object</li><li>• Protects other Thread methods (e.g. <code>start()</code>)</li></ul>        |
| CONS | <ul style="list-style-type: none"><li>• Cannot inherit from another class</li><li>• Danger of overriding other methods in Thread class (e.g. <code>start()</code>)</li></ul> | <ul style="list-style-type: none"><li>• Harder to access non-static Thread methods when defining Runnable objects</li><li>• Must create intermediate Runnable object in order to create Thread</li></ul> |

# Thread States

- What happens if we do the following?

```
Thread h1 = new HelloWorldThread ();
h1.start();
h1.start();
```

- Answer

```
Exception in thread "main"
java.lang.IllegalThreadStateException
```

- What?
  - Not every method is legal on every Thread object
  - The *state* of the object determines this validity
  - In this case, you cannot start a thread that has already been started

# Thread States?

- Accessible via method `Thread.State getState()`
- `Thread.State` is an enumerated type recording state of thread object
  - `NEW`  
A thread that has not yet started is in this state.
  - `RUNNABLE`  
A thread executing in the Java virtual machine is in this state.
  - `BLOCKED`  
A thread that is blocked waiting for a monitor lock is in this state.
  - `WAITING`  
A thread that is waiting indefinitely for another thread to perform a particular action is in this state.
  - `TIMED_WAITING`  
A thread that is waiting for another thread to perform an action for up to a specified waiting time is in this state.
  - `TERMINATED`  
A thread that has exited is in this state.

[Quoted from <http://docs.oracle.com/javase/6/docs/api/java/lang/Thread.State.html>]

# Thread State Example Revisited

```
Thread h1 = new HelloWorldThread (); // state is NEW
h1.start(); // state is RUNNABLE
h1.start(); // Error!
```

- When `h1` is created, its state is `NEW`
- After `h1.start()` is called, the state is `RUNNABLE`
- `h1.start()` can only be called when state is `NEW`!



# More on Thread States

- Some Thread methods (e.g. start) only applicable when object is in correct state
- The states NEW, RUNNABLE, TERMINATED are probably easiest to understand
- We will learn about the states BLOCKED, WAITING, TIMED\_WAITING later

# Other Thread State Methods

- `boolean isAlive()`
  - Returns `true` if thread has been started but is not terminated
  - `t.isAlive()` equivalent to  
`(t.getState() != NEW) && (t.getState() != TERMINATED)`
- `void join()`
  - Blocks until thread terminates, then terminates
  - `t.join()` very similar to  
`while (t.isAlive ()) { }`
- `void join(int millis)`

Like `t.join()` except that if `t` has not terminated in `millis` milliseconds, then `t.join(millis)` nevertheless terminates

# Threads and Process Termination

- A process (JVM) terminates when “there is nothing left that has to be done”
- When does this hold?
  - When the main thread terminates?
  - When all threads terminate?
  - When “the important” threads terminate?
- Java answer: when all *user threads* terminate

# User Threads vs. Daemon Threads

- In Java, every thread object is by default a *user thread*
- A Java process can terminate if and only if all user threads (including, but not only, main) have terminated
- Threads may be changed to *daemon threads* using method `setDaemon(boolean on)`
  - If the only nonterminated threads are daemons, then the JVM will terminate
  - Daemon threads should only be used for “background work” (e.g. updating status bars, etc.) needed while “useful” computation is being performed
- `setDaemon()` is only valid if thread state is `NEW`; otherwise, `IllegalThreadStateException` thrown

# More on Thread Termination

- When a thread object terminates, the object still remains!
  - Thread state is `TERMINATED` ...
  - ... but object still exists


# Thread Execution

- Once threads are started, what determines when they are eligible for execution?
- Answer: scheduler!
  - Operating system routine responsible for allocating processor time to threads
  - If there are more processors than threads, could allocate each thread to its own processor
  - If there are more threads than processors, *time-slicing* may be needed to *interleave* access to processors
    - Each thread executes for a while, then is pre-empted
    - Exact scheme also takes account of priorities, also whether or not threads are blocked
    - What if thread is in the middle of something “atomic”?


# Scheduling Example (1)



# Scheduling Example (2)



*Threads of a process allowed to run on either CPU*



*p2 threads:*   *p1 threads:*  



# Methods for Interacting with Scheduler

- `void setPriority(int newPriority)`  
Set priority to given value (must be between MIN\_PRIORITY and MAX\_PRIORITY: see below)
- `int getPriority()`  
Return priority value
- `static void yield()`  
“Hint” to scheduler that thread can give up processor
- `static void sleep(int millis)`  
Block for millis milliseconds
- `static int MIN_PRIORITY`  
Smallest (lowest) priority
- `static int MAX_PRIORITY`  
Largest (highest) priority
- `static int NORM_PRIORITY`  
Default priority

# InterruptedException

- Thrown by some Thread methods (e.g. `sleep( )`)
  - Raised when a thread is interrupted while sleeping
  - We will see about interruptions later
- When you call such a method, you must either
  - Catch the exception, e.g.

```
try { ... sleep (1000); ... }
catch (InterruptedException e) { ... }
```
  - ... or include a `throws` directive in your method declaration, e.g.

```
public void myMethod (...) throws
 InterruptedException {...}
```

# currentThread()

```
static Thread currentThread()
```

- Returns thread of current execution
- Useful when implementing thread operations via Runnable, as you can get access to thread info at runtime